


EVML

ARTIFICIAL NEURAL NETWORKS

JEROEN VEEN


ADVERSARIAL SWEATER OF DOOM



[Render Yourself Invisible To AI With This Adversarial Sweater Of Doom | Hackaday](#), 20 October 2022

AGENDA

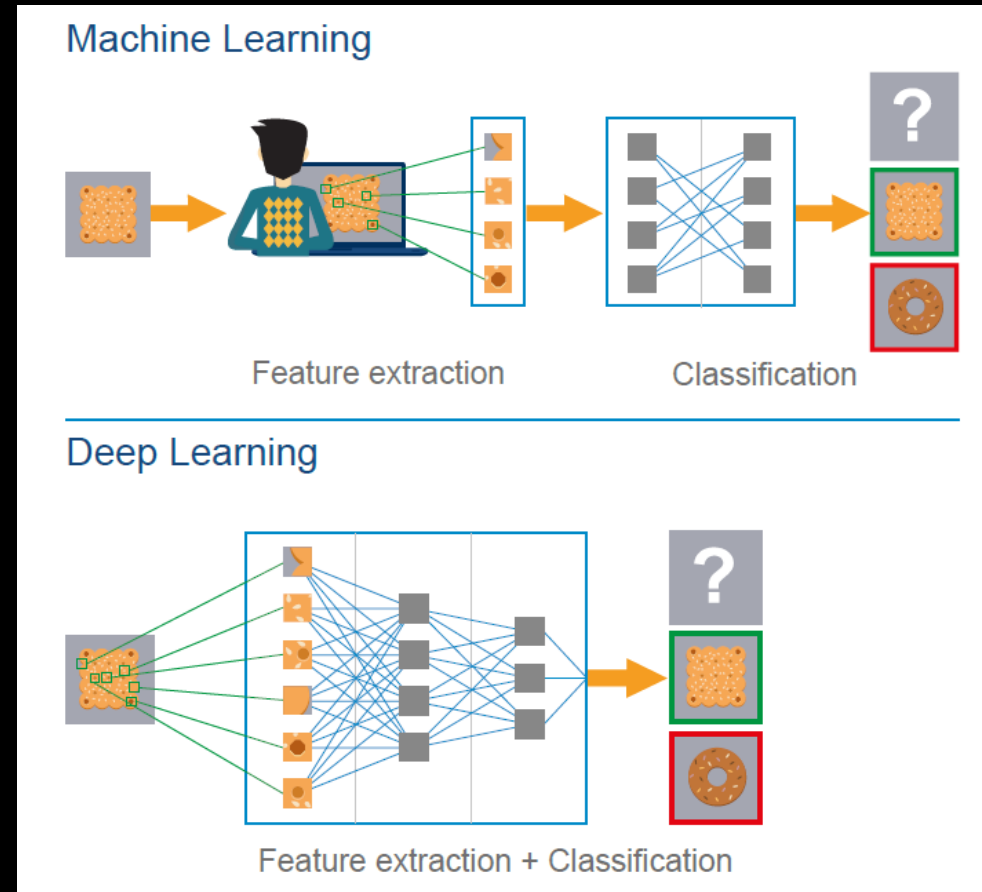
- Machine learning vs deep learning
- Biological neuron
- Perceptron
- Multi-layer perceptron (MLP)
- Backpropagation
- Regression and classification MLP



BACKGROUND MATERIAL


- https://deeplizard.com/learn/playlist/PLZbbT5o_s2xq7Lwl2y8_QtvuXZedL6tQU
- <https://www.3blue1brown.com/topics/neural-networks>
- [MIT Deep Learning 6.S191 \(introtodeeplearning.com\)](https://introtodeeplearning.com/)

MACHINE LEARNING VS DEEP LEARNING



Autonomous
feature
definition

DEFINING AI, DL & ML



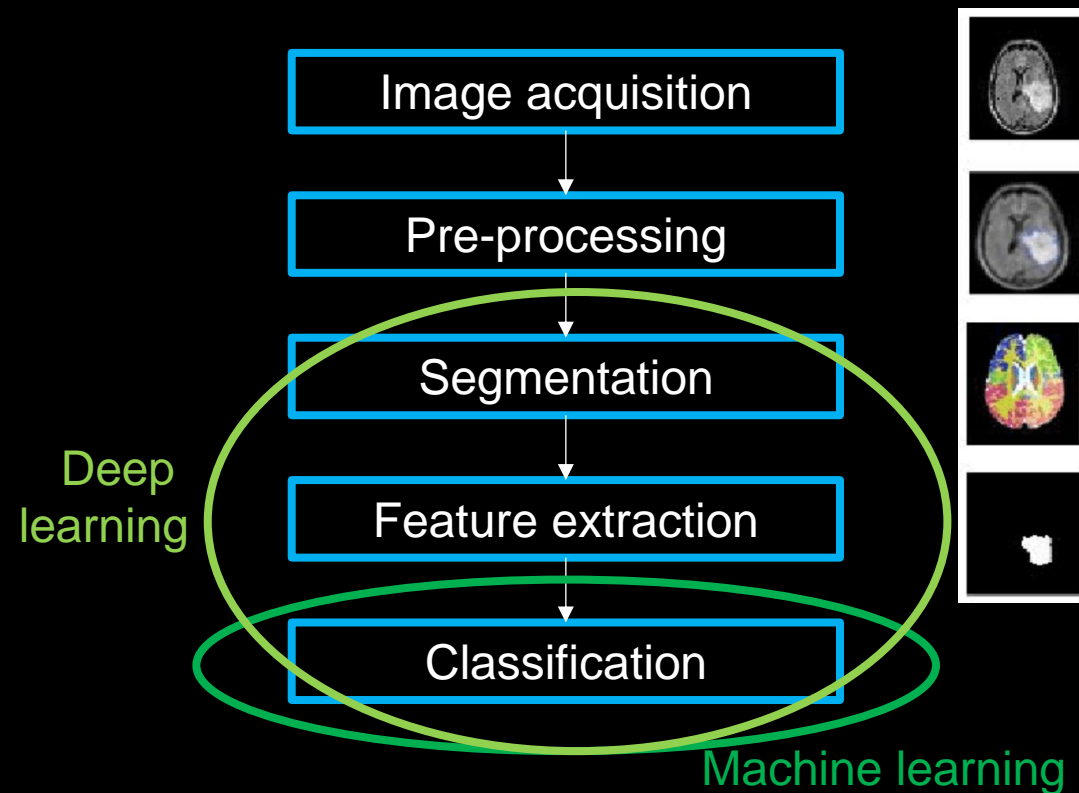
- Strong AI vs Applied AI
- Cognitive replication
- Rational process

Machine learning

- Performs predictive analysis
- Just fancy math & pattern matching



MACHINE LEARNING APPLIED TO VISION

- Classical image processing



APPLICATION AREAS OF DEEP LEARNING

- Anomaly detection, image classification, image segmentation and object recognition.
- Higher precision and greater flexibility compared to conventional image analysis methods.




Source: Tilmann Zuper, Artificial Intelligence in Image Processing: Deep Learning Compared with Conventional Methods, Basler AG, Ahrensburg, Germany.

COSTS OF DEEP LEARNING

- Additional hardware
Large memory and computing capacity is required, typically outsourced to e.g. GPUs (graphic cards).
- Power consumption:
Large memory and computing capacity increase power consumption and thus the heat generation. This can be problematic for embedded systems.
- High amount of training data:
Large number of training images required, which is sometimes difficult in the development of a Machine Vision application.

ON THE EDGE




Source: <https://medium.com/da-labs/edge-ai-the-future-of-ai-d954ebc40a46>

HYBRID APPROACH


- High performance with low memory and power requirement
- Image preprocessing with conventional methods.
- An artificial neural network then delivers the desired results with the preprocessed data.
- DL mingled with expert systems

BIOLOGICAL NEURONS




NEURAL CIRCUITS

- Population of neurons interconnected by synapses to carry out a specific function when activated
- Highly complex computations can be performed by a network of fairly simple neurons



Source: Géron, ISBN: 9781492032632


BRAINBOW OF CEREBRAL CORTEX NEURONS LABELED WITH DIFFERENT COLORS



Source: Nature, Meet Nurture - Neuroscience News credited to Lichtman Lab, Harvard University.

THRESHOLD LOGIC UNIT (TLU)


- Elementary unit of an ANN
- Simplified model of a biological neuron
- Dot product followed by a non-linear function
- Performs linear binary classification



Source: Géron, ISBN: 9781492032632

PERCEPTRON

- Single layer of TLUs
- Multioutput classifier
- Connection weights




Source: Géron, ISBN: 9781492032632

OUTPUT COMPUTATION

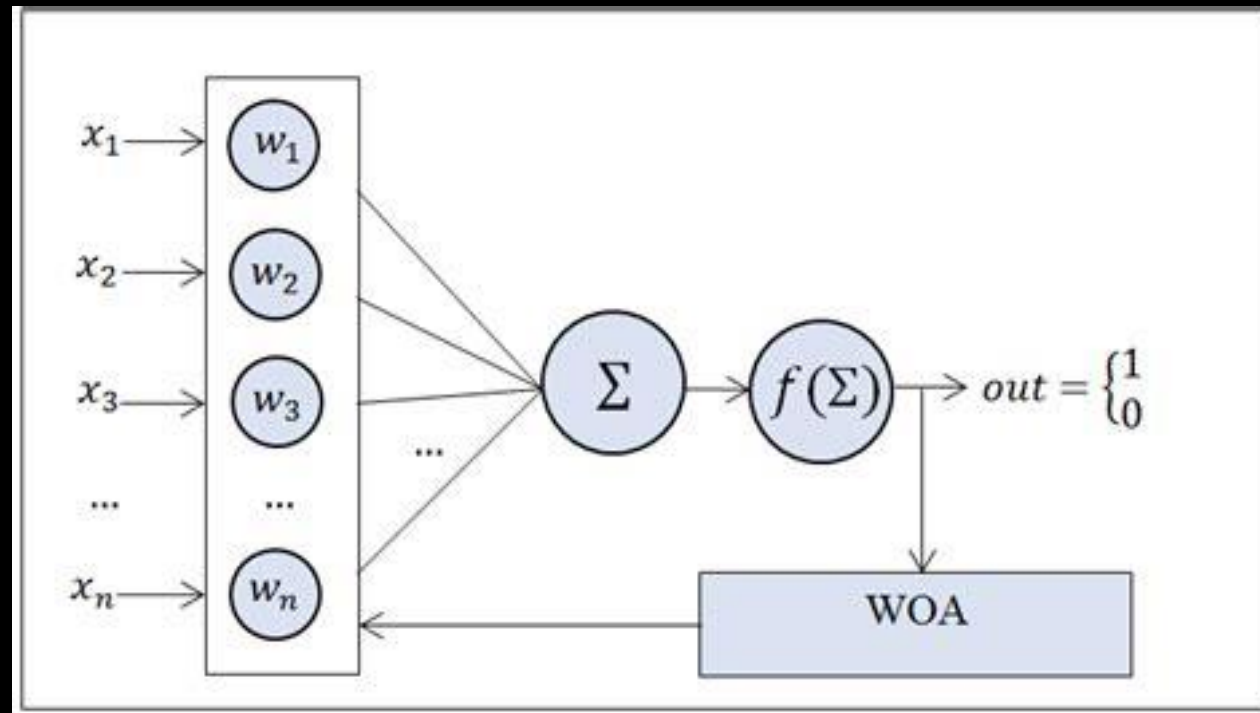
$$\underline{y}(X) = \phi(XW + \underline{b})$$

Output vector Matrix of input features Activation function Weight matrix Bias vector



HOW TO FIND THE OPTIMAL WEIGHTS?

- Optimization
- Cost function




PERCEPTRON TRAINING ALGORITHM

- Multi-dimensional optimization problem
- Gradient descent

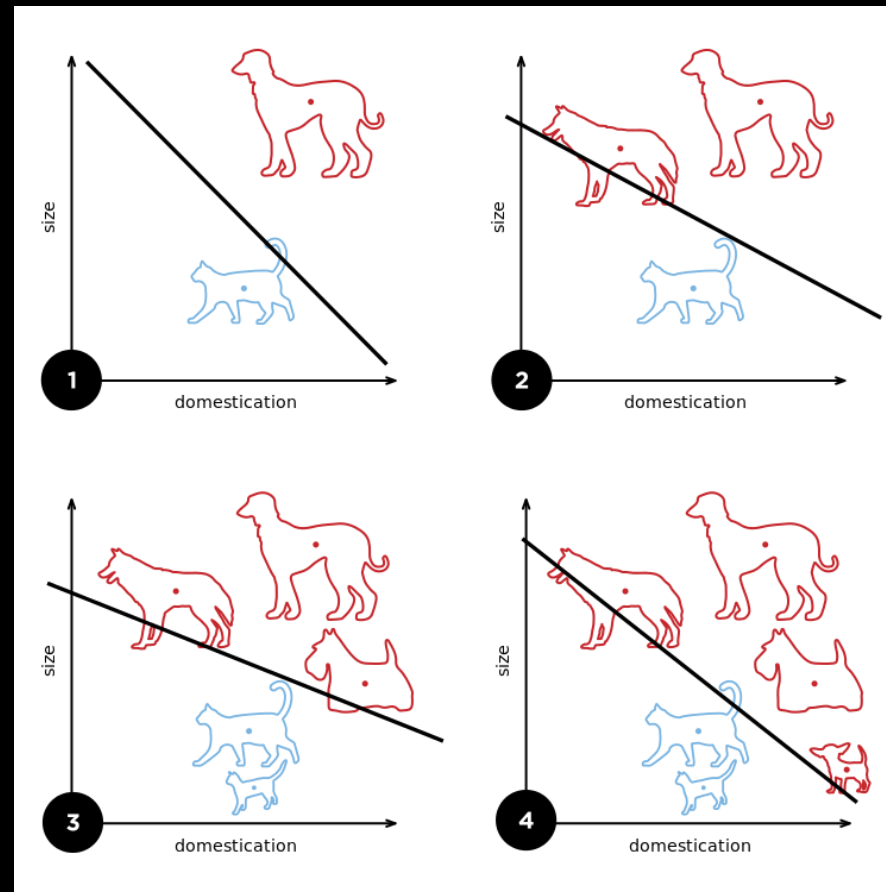
$$w_{i,j}^{(\text{next step})} = w_{i,j} + \eta (y_j - \hat{y}_j) x_i$$

Diagram illustrating the weight update formula for the Perceptron Training Algorithm. The formula is shown with annotations:

- $w_{i,j}$ is labeled "Connection weights".
- η is labeled "Learning rate".
- $(y_j - \hat{y}_j)$ is labeled "error".
- x_i is labeled "Input value".




EXAMPLE OF ITERATIVE UPDATING



Source: <https://en.wikipedia.org/wiki/Perceptron>

PERCEPTRON LIMITATIONS


- Linear decision boundary
- Incapable of learning complex patterns



Source: <https://mc.ai/solving-xor-with-a-single-perceptron/>

MULTILAYER PERCEPTRON

- Feedforward neural network



Source: Géron, ISBN: 9781492032632

BACKPROPAGATION

Let's now watch

MIT's intro to deep learning

<https://www.youtube.com/watch?v=7sB052Pz0sQ?t=35m38s>

3BLUE1BROWN SERIES S3 • A3

What is backpropagation really doing?


BACKPROPAGATION

- for each training instance, the backpropagation algorithm first makes a prediction (forward pass) and measures the error,
- then goes through each layer in reverse to measure the error contribution from each connection (reverse pass),
- and finally tweaks the connection weights to reduce the error (Gradient Descent step).

ACTIVATION FUNCTIONS


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout


$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



CLASSIFICATION MLP



Source: <http://rinterested.github.io/statistics/softmax.html>

REGRESSION MLP

- No activation function for output neurons required
- Use functions to bound outputs, e.g. relu, softplus, logistic function

Table 10-1 summarizes the typical architecture of a regression MLP.

Table 10-1. Typical regression MLP architecture

| Hyperparameter | Typical value |
|----------------------------|--|
| # input neurons | One per input feature (e.g., $28 \times 28 = 784$ for MNIST) |
| # hidden layers | Depends on the problem, but typically 1 to 5 |
| # neurons per hidden layer | Depends on the problem, but typically 10 to 100 |
| # output neurons | 1 per prediction dimension |
| Hidden activation | ReLU (or SELU, see Chapter 11) |
| Output activation | None, or ReLU/softplus (if positive outputs) or logistic/tanh (if bounded outputs) |
| Loss function | MSE or MAE/Huber (if outliers) |

Source: Géron, ISBN: 9781492032632

EXERCISE

- <https://developers.google.com/machine-learning/crash-course/reducing-loss/playground-exercise>
- How did the lower learning rate impact convergence?
- Can you find a learning rate too slow to be useful?
- Better website:
<https://playground.tensorflow.org>