

EVML


DEEP NEURAL NETWORKS

JEROEN VEEN

CONTENTS

- Recap ANN and example
- Vanishing and exploding gradients
- Transfer learning
- Training optimization
- Learning rate scheduling
- Regularization


Use: <https://alexlenail.me/NN-SVG/index.html>
to draw nice maps



AGENDA

- DL portfolio walkthrough
- Tensorboard, visualizing the training process
- Data augmentation
- Storing and loading models
- Experimental route to optimize architectures
- Fine-tuning neural network hyperparameters

RECALL ANNS




Source: Géron, ISBN: 9781492032632

ACTIVATION FUNCTIONS


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU


$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



Source: http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture6.pdf

TRAINING

- Multi-dimensional optimization problem
- Gradient descent



$$w_{i,j} \text{ (next step)} = w_{i,j} + \eta (y_j - \hat{y}_j) x_i$$

Diagram illustrating the weight update formula for gradient descent:

- $w_{i,j}$ (next step): Connection weights
- η : Learning rate
- $(y_j - \hat{y}_j)$: error
- x_i : Input value

EXERCISE: BASIC IMAGE CLASSIFICATION

- Train a shallow net, see also Géron 10.2.1
- Trouble downloading the datasets? Let me know
- Training can take quite some time...
 - >> Use a small number of weights in the hidden layers
 - >> Work on colab.research.google.com
- Train and validate


See [tf_quickstart.py](#)



Instead of passing a validation set using the `validation_data` argument, you could set `validation_split` to the ratio of the training set that you want Keras to use for validation. For example, `validation_split=0.1` tells Keras to use the last 10% of the data (before shuffling) for validation.

UNSTABLE GRADIENTS

- https://www.youtube.com/watch?v=qO_NLVjD6zE&t=105s
- Variance of the outputs > variance inputs leads to saturation



SOLUTIONS TO SPEED UP TRAINING

- Unstable gradients results in very slow training
- Smart weight initialization
- Non-saturating activation function
- Batch normalization

GLOROT AND HE INITIALIZATION

- Signals need to flow properly in both directions
- fan_{in} = number of inputs
- fan_{out} = number of neurons

Table 11-1. Initialization parameters for each type of activation function


Initialization	Activation functions	σ^2 (Normal)
Glorot	None, tanh, logistic, softmax	$1 / \text{fan}_{\text{avg}}$
He	ReLU and variants	$2 / \text{fan}_{\text{in}}$
LeCun	SELU	$1 / \text{fan}_{\text{in}}$

- See also <https://www.youtube.com/watch?v=8krd5qKVw-Q&t=300>

Fan-Out and Fan-In


Fan-out – number of load gates connected to the output of the driving gate

- gates with large fan-out are slower





Fan-in – the number of inputs to the gate

- gates with large fan-in are bigger and slower



NONSATURATING ACTIVATION FUNCTIONS


- Standard ReLU may lead to 'dying' neurons
- ReLU variants: leaky ReLU, PReLU, ELU, SELU



Source: Géron, ISBN: 9781492032632

BATCH NORMALIZATION


- Mitigate both vanishing and exploding gradients
- Normalize each hidden layer's input during training



- Input mean and std (μ, σ), output scale and offset (γ, β) parameters are learned over entire batch

TRANSFER LEARNING

- Reuse pretrained (lower) layers to speed up training, requiring significantly less data
- Works best when the inputs have similar low-level features



gradual
unfreezing
from the top

iterate until
right number
of layers

Source: Géron, ISBN: 9781492032632


TRAINING OPTIMIZATION

- Cost functions / Loss functions
measure of correctness of a prediction
- e.g. mean squared error
cross entropy
log loss
- Descending the error curve, feedback on error
- Different kinds of optimization : gradient descent, stochastic gradient descent, adagrad, adam, etc.

a measure of dissimilarity
between the ground truth label
probability and the predicted
probability of the label

FASTER OPTIMIZERS

- Adaptive learning rate algorithms
- <https://www.youtube.com/watch?v=mdKjMPmcWjY&t=133s>




Source: Géron, ISBN: 9781492032632

OPTIMIZER COMPARISON


Class	Convergence speed	Convergence quality
SGD	*	***
SGD(momentum=...)	**	***
SGD(momentum=..., nesterov=True)	**	***
Adagrad	***	* (stops too early)
RMSprop	***	** or ***
Adam	***	** or ***
Nadam	***	** or ***
AdaMax	***	** or ***

OPTIMIZING CONSTANT LEARNING RATE

- Compare learning curves for various rates and pick optimal, but constant rate




Source: Geron, ISBN: 9781492032632



Source: <https://www.jeremyjordan.me/nn-learning-rate/>

LEARNING RATE SCHEDULING

- Strategies to adapt learning rate while training progresses.
- Power, exponential, piecewise constant scheduling:
drop learning rate every iteration, e.g.
- Performance scheduling:
reduce learning rate based on error
- (1cycle) scheduling



REGULARIZATION

- ℓ_1 regularization: sparse model
- ℓ_2 regularization: constrain weights
- (MC) dropout: ignore neurons during training
- Max-Norm regularization

L_2 and L_1 penalize weights differently:

- L_2 penalizes $weight^2$.
- L_1 penalizes $|weight|$.

