EVML

ADVANCED CNN

JEROEN VEEN

HAN_UNIVERSITY
OF APPLIED SCIENCES

DL REPORT TEMPLATE

- Questions?

AGENDA

- Object detection
- Object tracking
- Semantic segmentation
- Variational autoencoder

TYPICAL CNN ARCHITECTURE

- Perform classification
- Conv, pool, dense layers

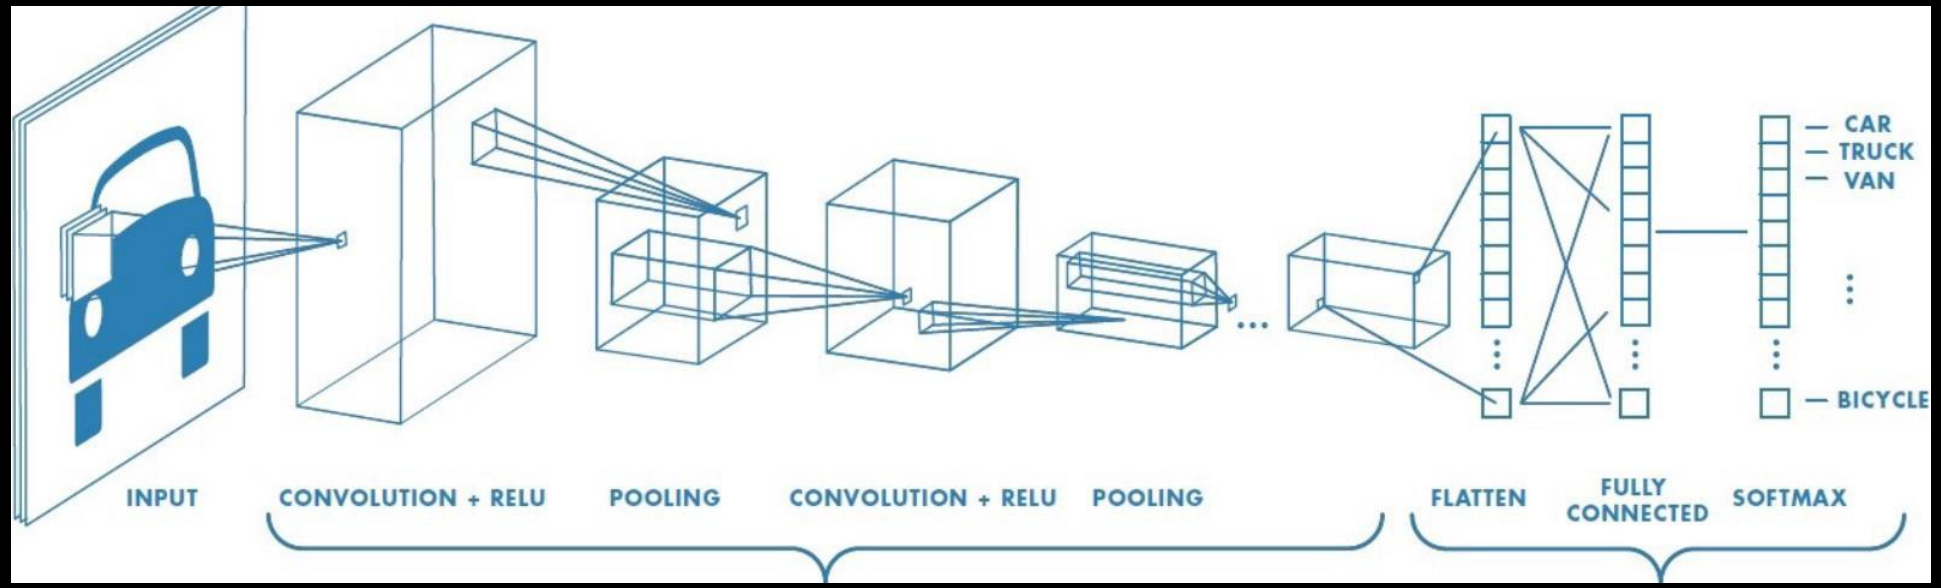

Source: https://towardsdatascience.com/a-comprehensive-guide-to-convolutional-neural-networks-the-eli5-way-3bd2b1164a53

CNN SUMMARY

- Emulate the behavior of a visual cortex (e.g. receptive fields)
- Higher-level representations of image content
- No feature definition, but automated extraction

- Biologically inspired perceptrons
- Multilayer perceptrons usually mean fully connected networks, which makes them prone to overfitting
- CNNs can be considered as regularized versions of multilayer perceptrons

CONVOLUTIONAL LAYER SUMMARY

- Local connectivity
- Shared weights
- 3D volumes of neurons
- Output is a stack if feature maps

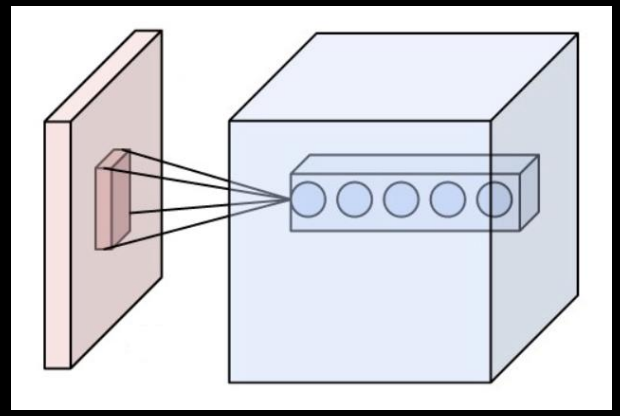

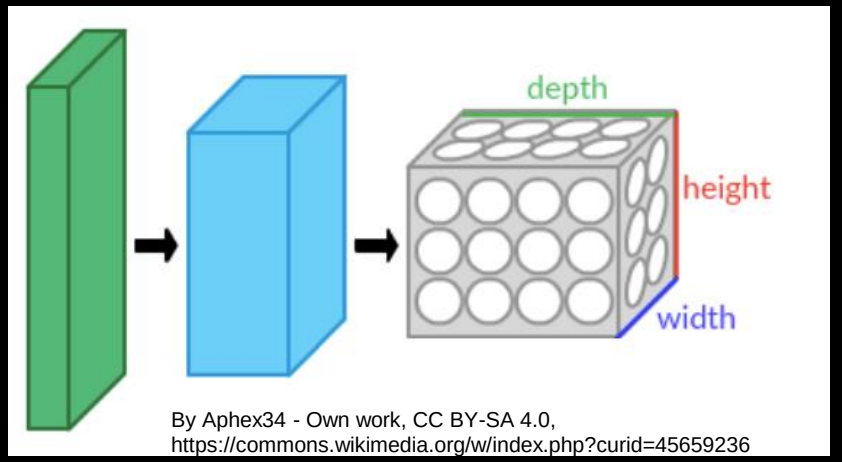

By Aphex34 - Own work, CC BY-SA 4.0, https://commons.wikimedia.org/w/index.php?curid=45659236

MORE TRANSFER LEARNING HINTS

- https://keras.io/api/applications

CLASSIFICATION AND LOCALIZATION

- Add second dense output layer to predict coordinates (regression)

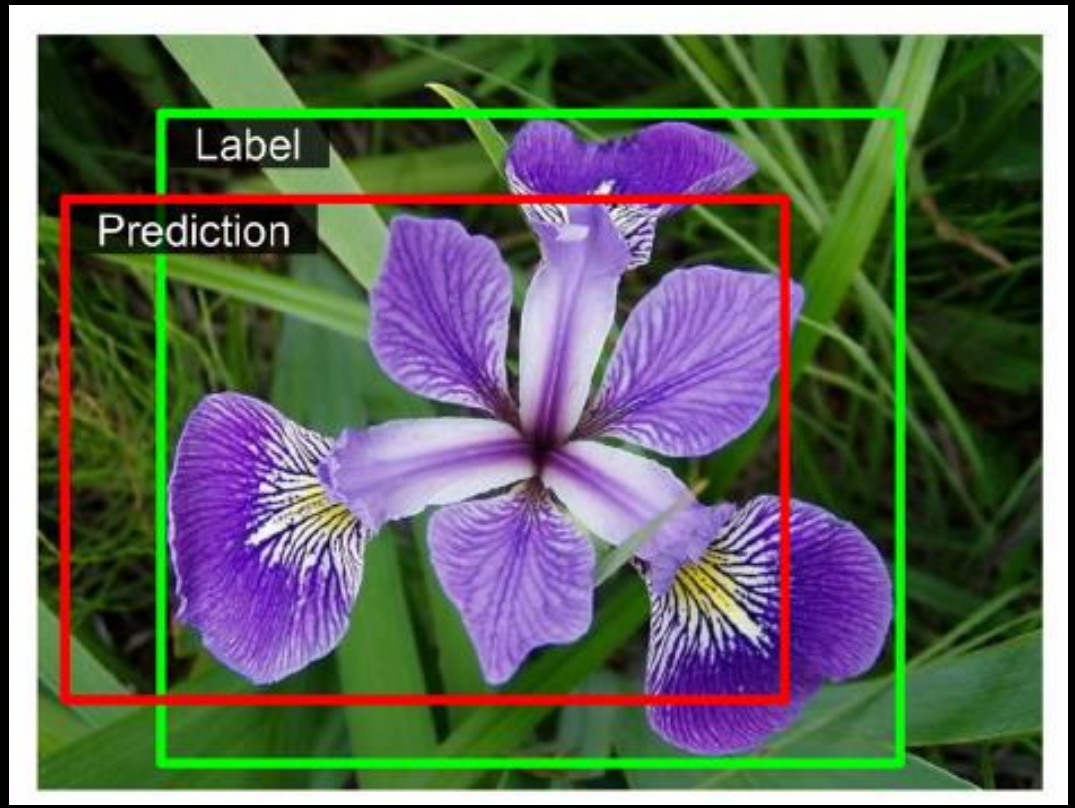

Source: Géron, ISBN: 9781492032632

JACCARD INDEX

- Performance metric: intersection over Union (IoU)
- measures similarity between finite sample sets

$$IoU = \frac{\text{Area of Overlap}}{\text{Area of Union}}$$

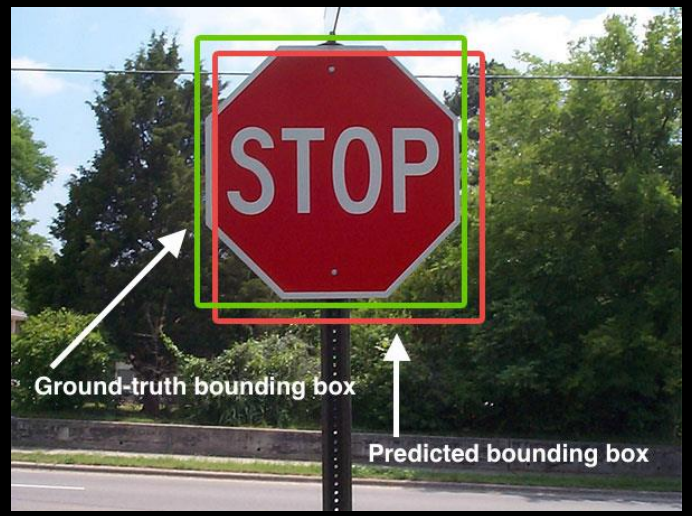

Source: By Adrian Rosebrock - http://www.pyimagesearch.com/2016/11/07/intersection-over-union-iou-for-object-detection/, CC BY-SA 4.0, https://commons.wikimedia.org/w/index.php?curid=57718560

OBJECT DETECTION

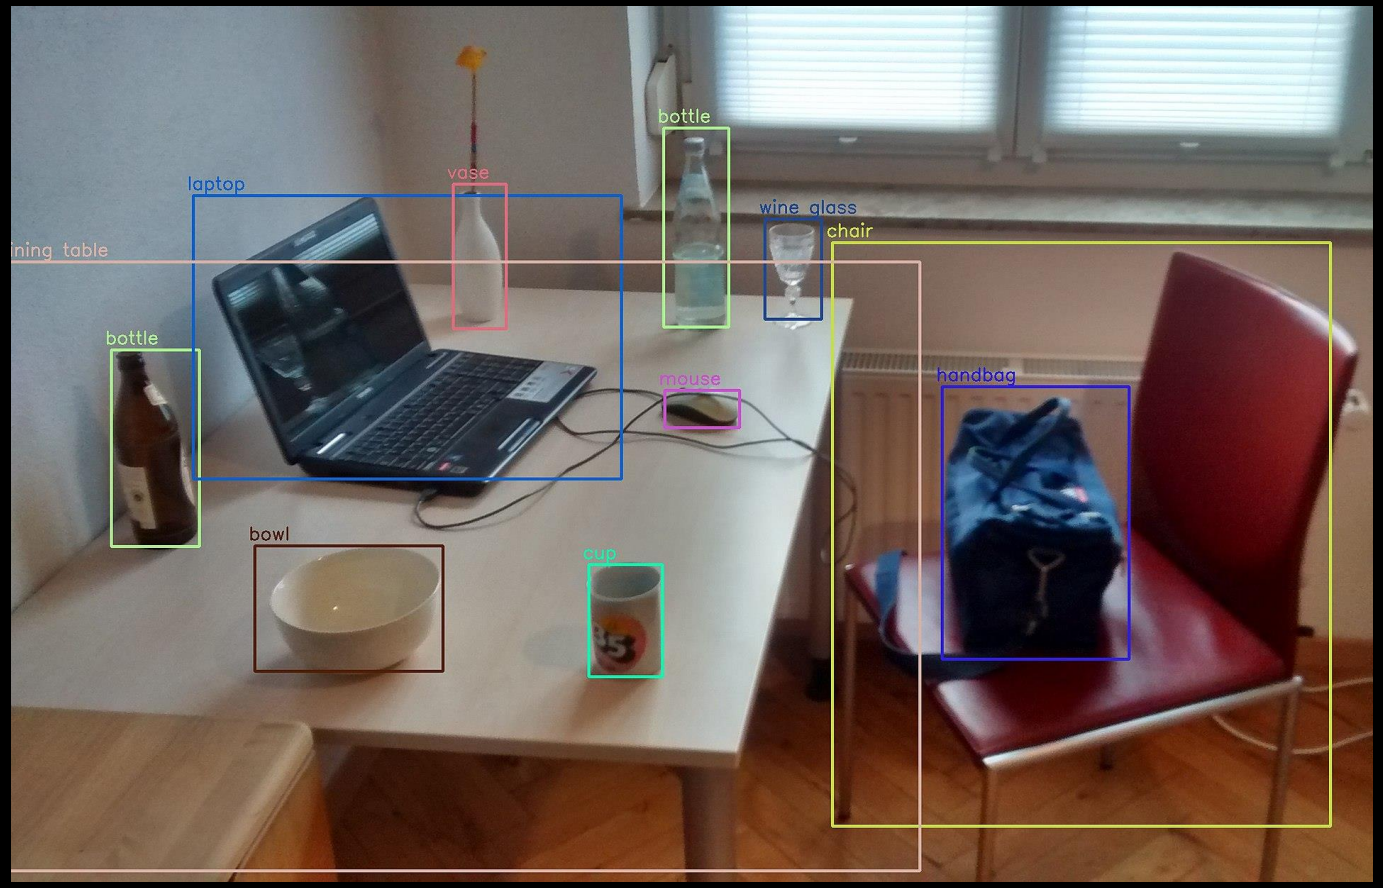

By (MTheiler) - Own work, CC BY-SA 4.0, https://commons.wikimedia.org/w/index.php?curid=75843378

SLIDING WINDOW

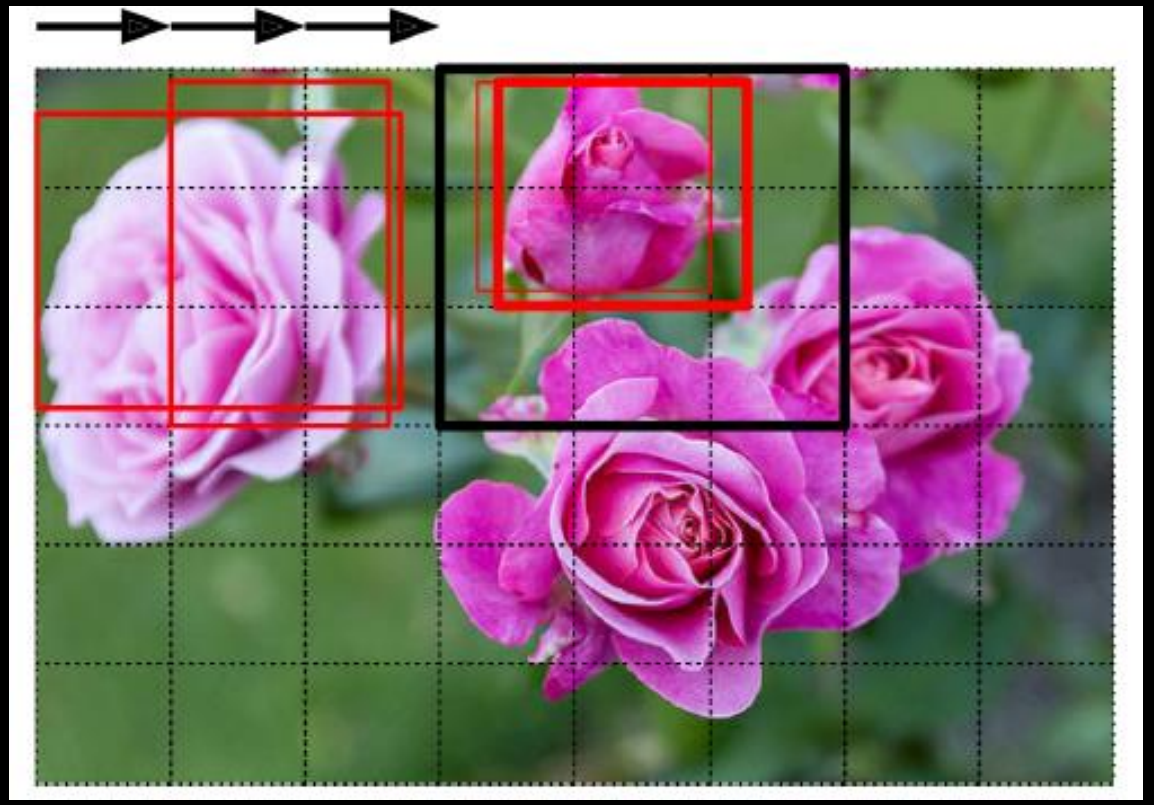

Source: Géron, ISBN: 9781492032632

SLIDING WINDOWS

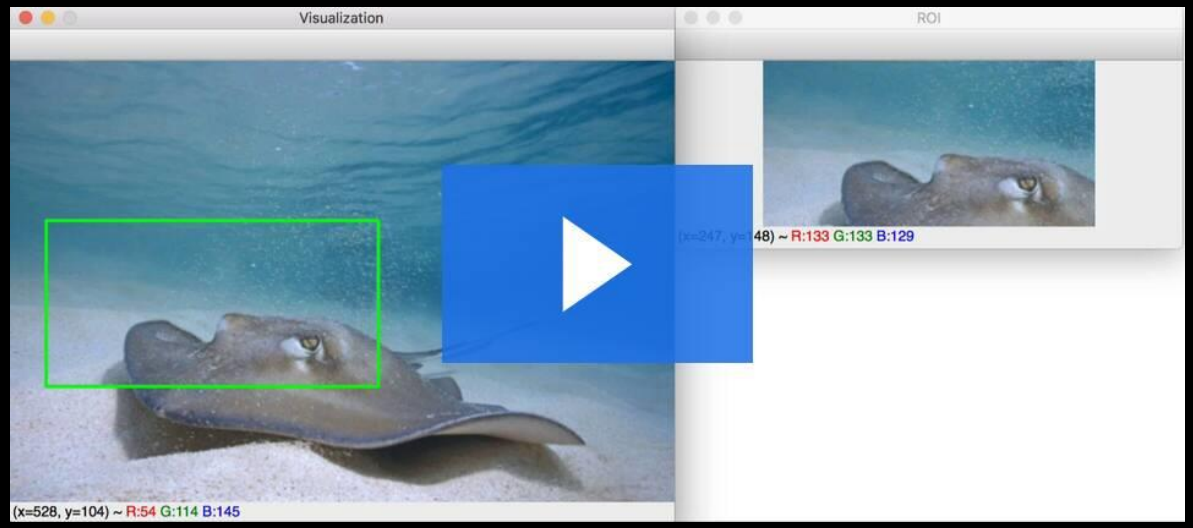

Source: https://www.pyimagesearch.com/2020/06/22/turning-any-cnn-image-classifier-into-an-object-detector-with-keras-tensorflow-and-opencv/

REGION-BASED CNN

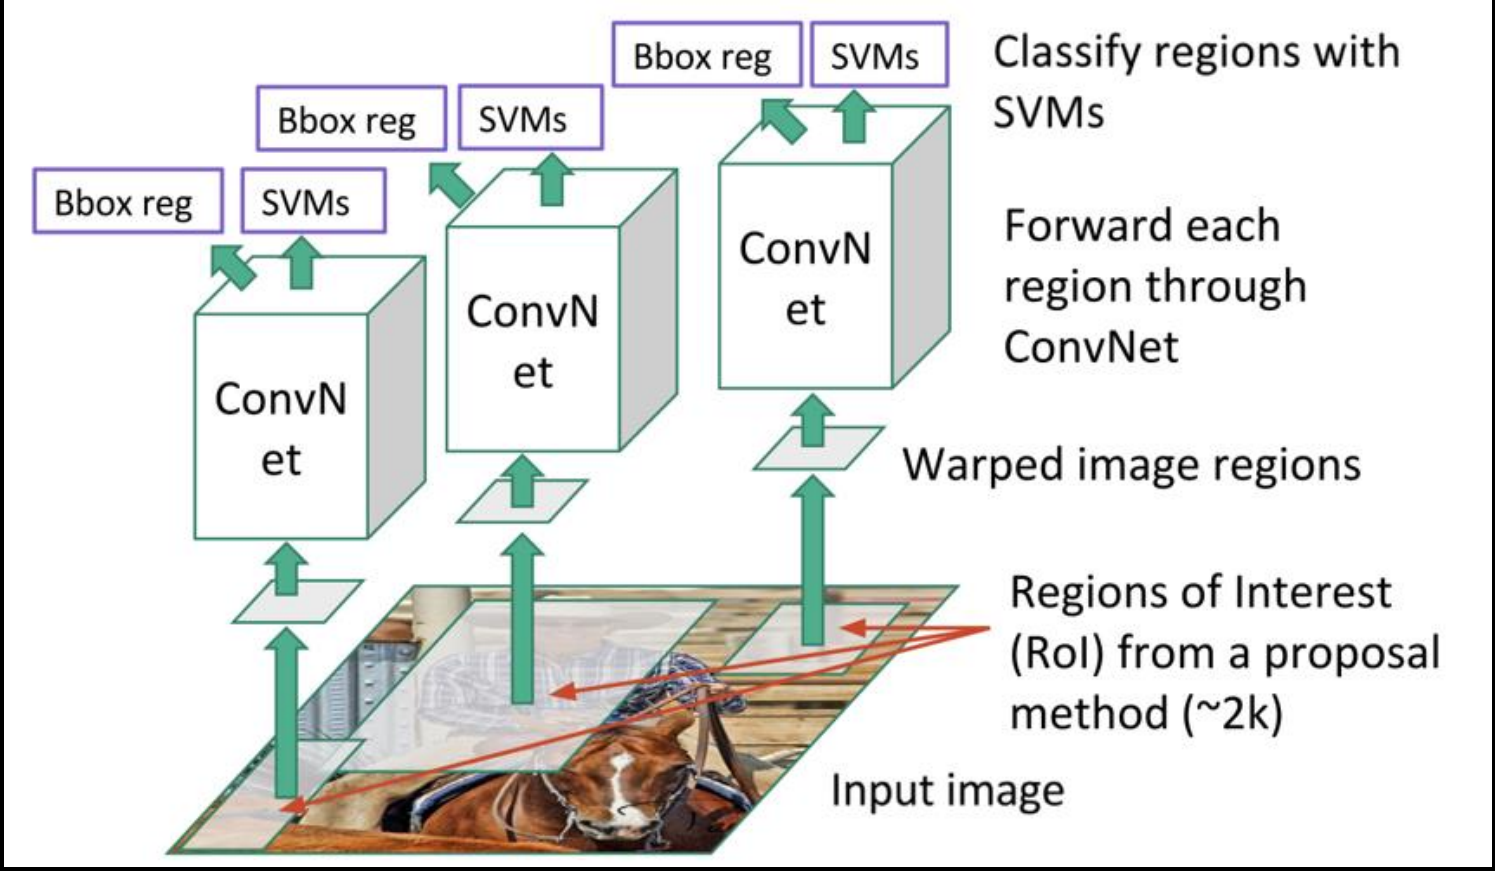

FPN: FEATURE PYRAMID NETWORK

- FPN: Feature Pyramid Network (2016) - KiKaBeN

Faster R-CNN

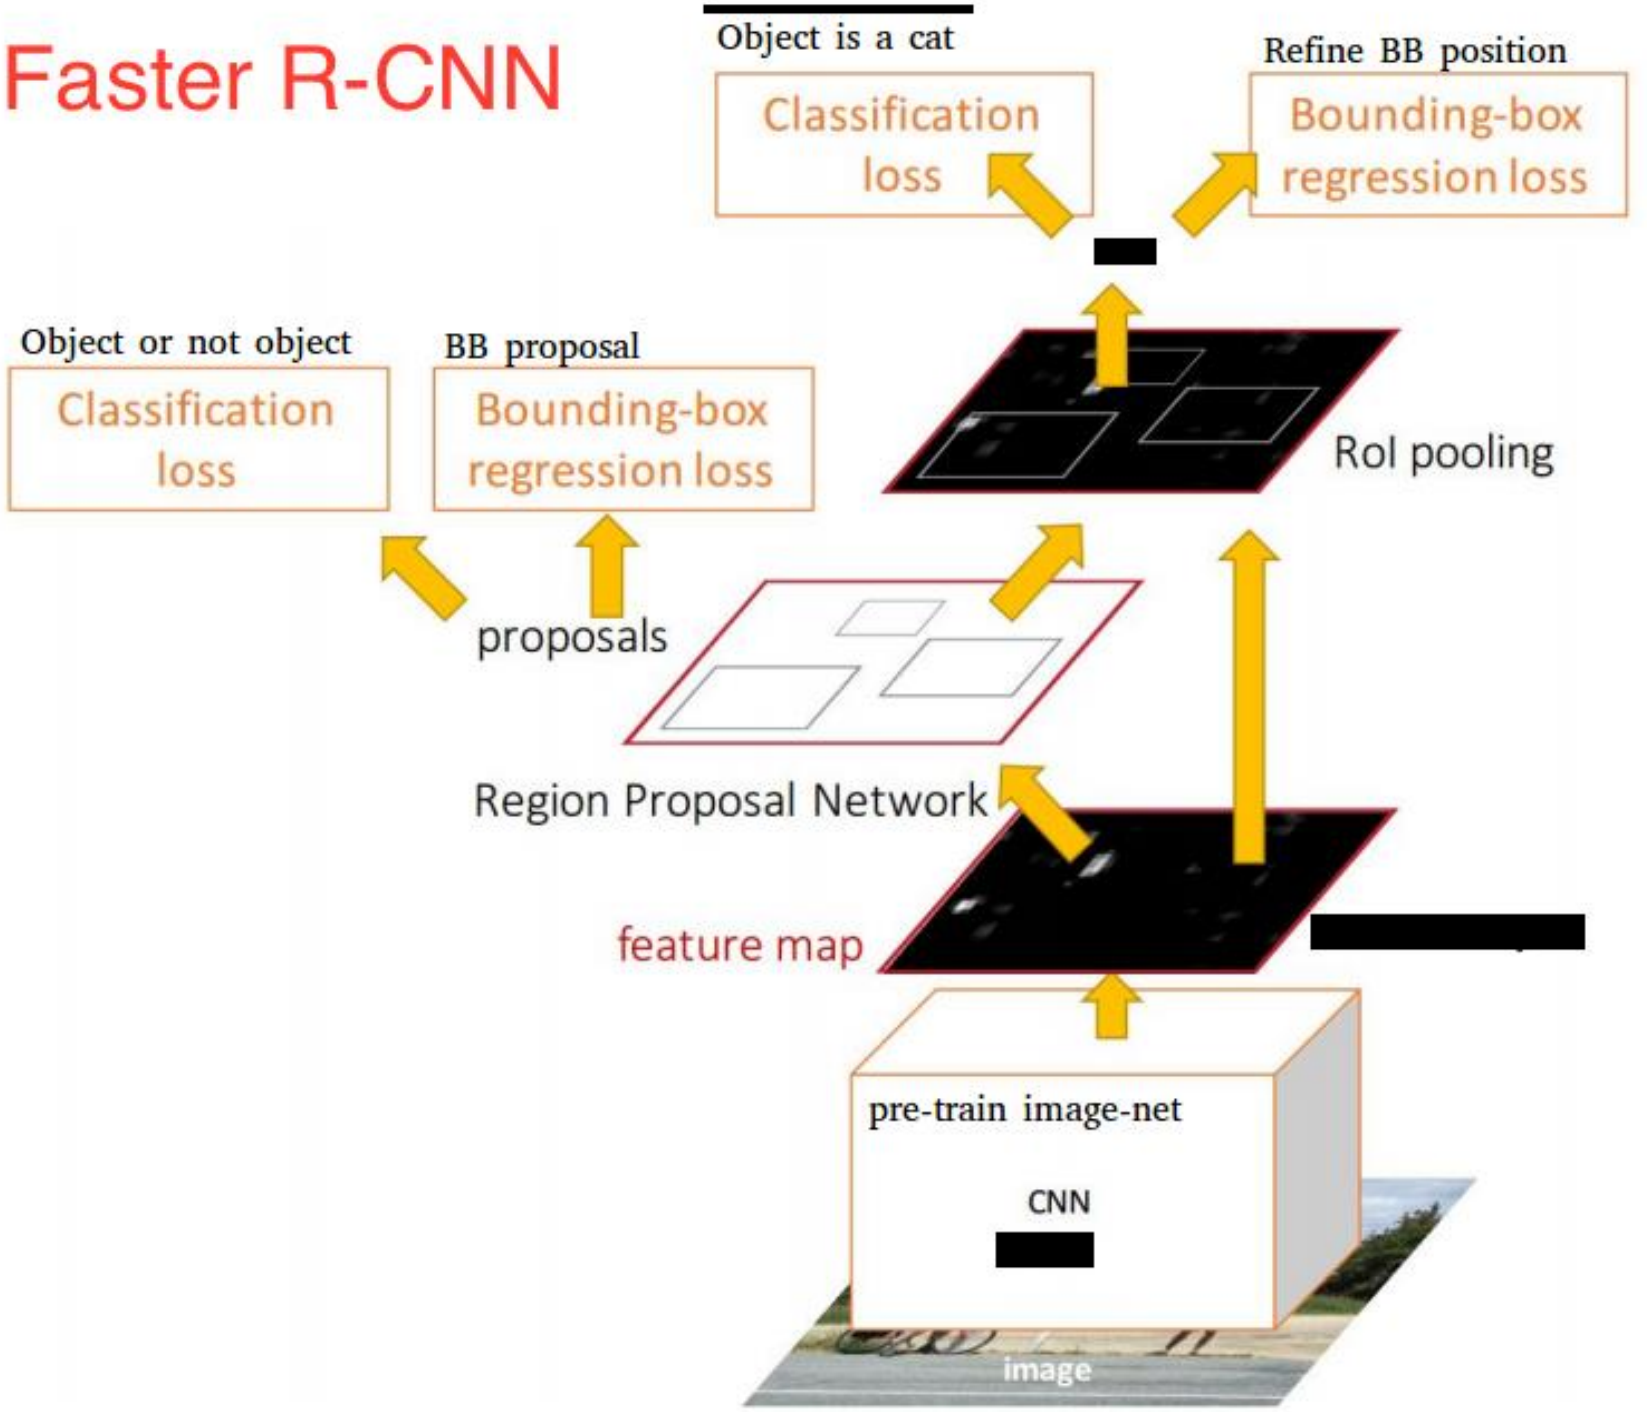
Object is a cat
Classification loss
Refine BB position
Bounding-box regression loss
Object or not object
Classification loss
BB proposal
Bounding-box regression loss
RoI pooling
proposals
Region Proposal Network
feature map
pre-train image-net
CNN
image

OBJECT DETECTION METHODS

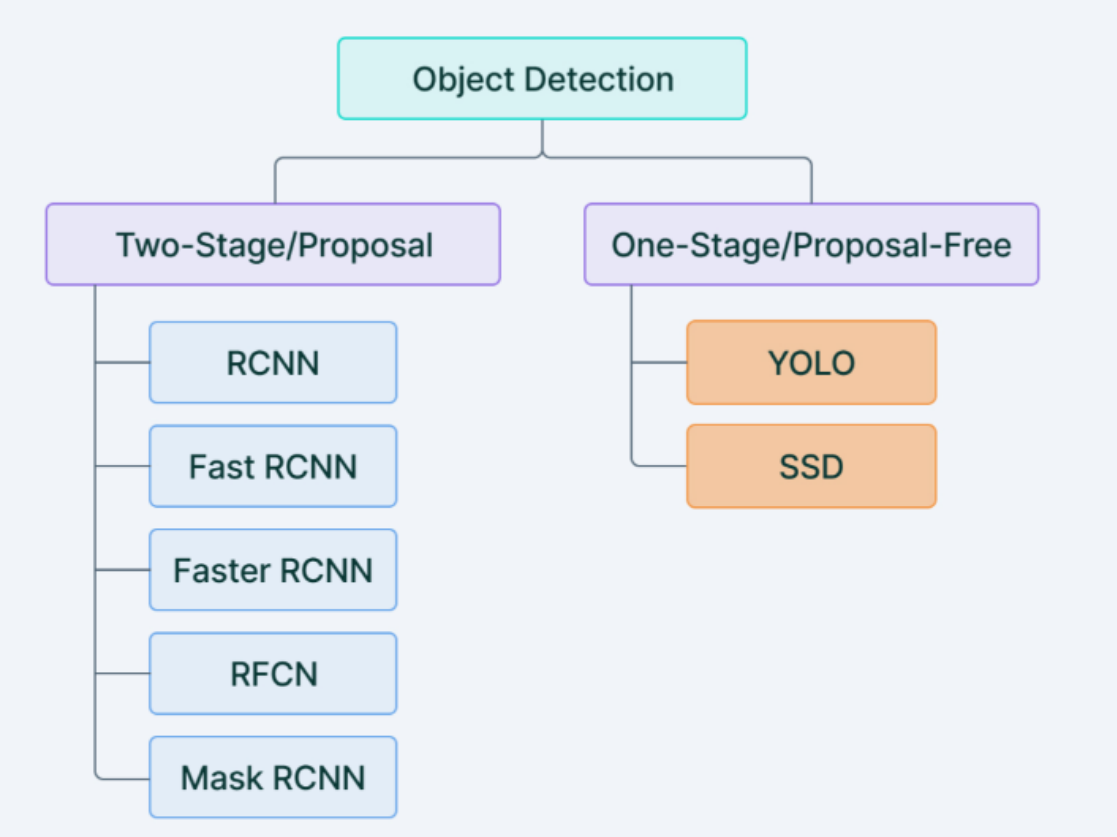

YOLO V9

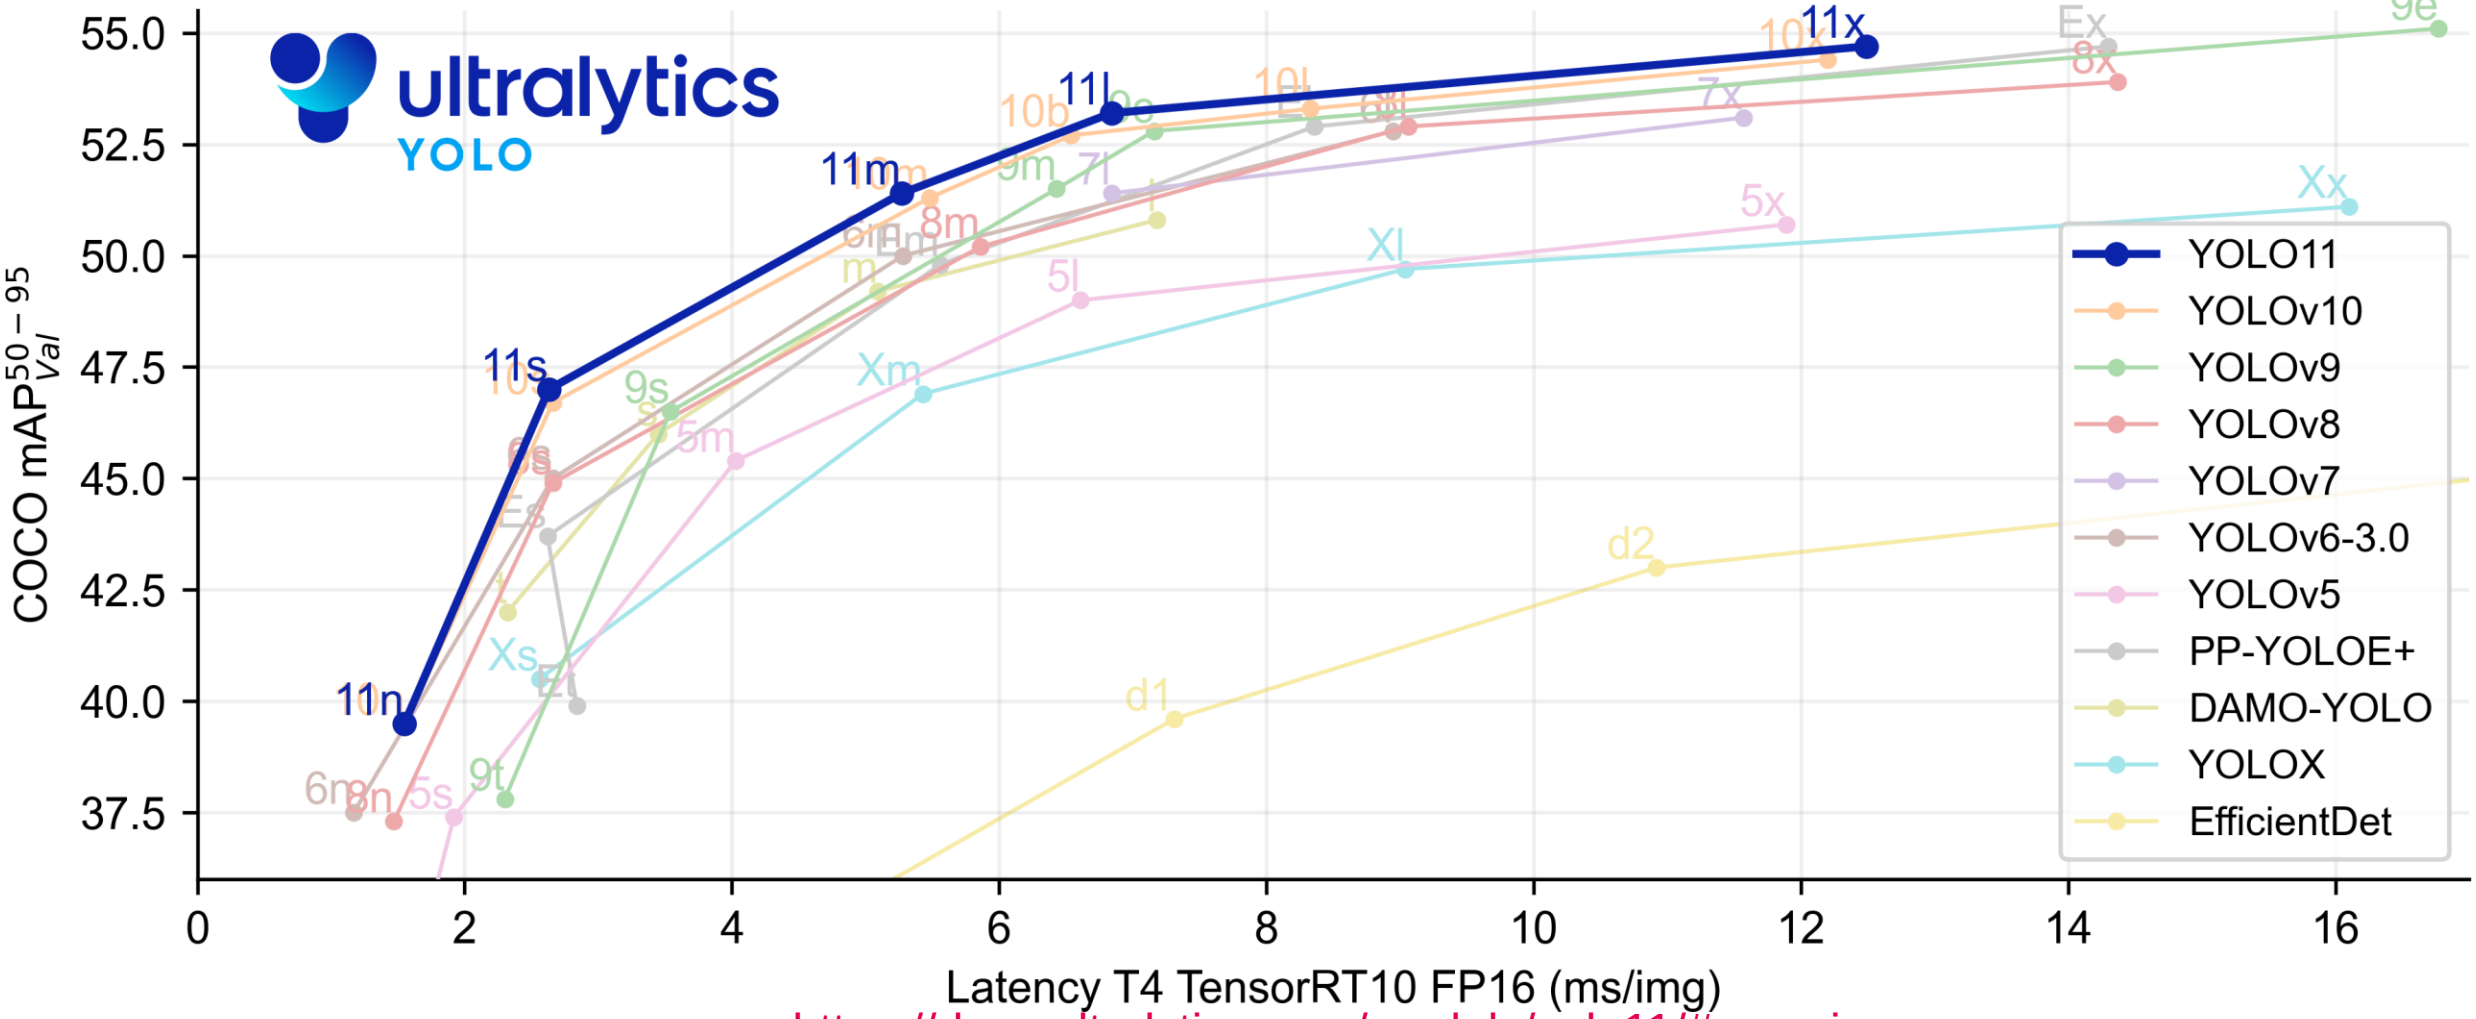

https://docs.ultralytics.com/models/yolo11/#overview

YOLO

- Intro: https://youtu.be/LvArU0AH8s8?t=35s
- Explanation: https://www.youtube.com/watch?v=svn9-xV7wjk&t=170s

  I'll give a short version ☺

YOU ONLY LOOK ONCE (YOLO)

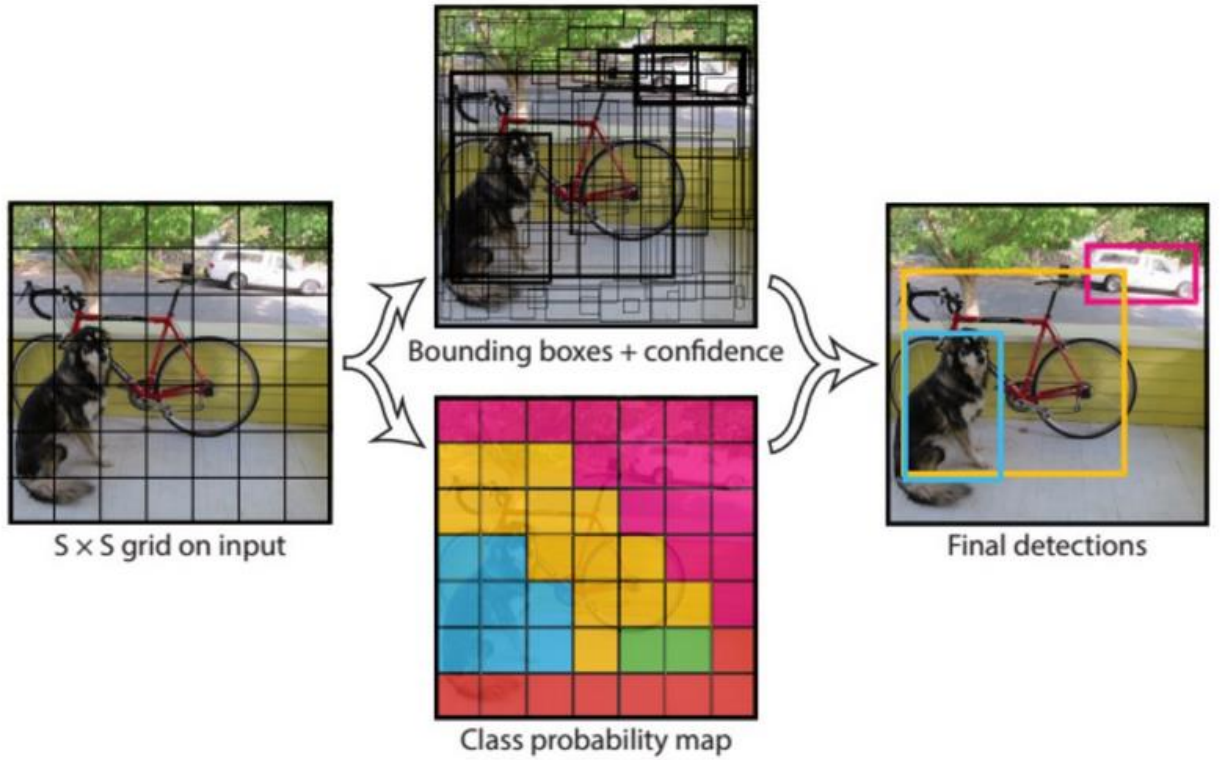

Source: Redmon et al., You Only Look Once:Unified, Real-Time Object Detection, 2016.

See: https://youtu.be/NM6lrxy0bxs?list=PLrrmP4uhN47Y-hWs7DVfCmLwUACRigYyT
https://pjreddie.com/darknet/yolo/

YOLO BASIC ARCHITECTURE

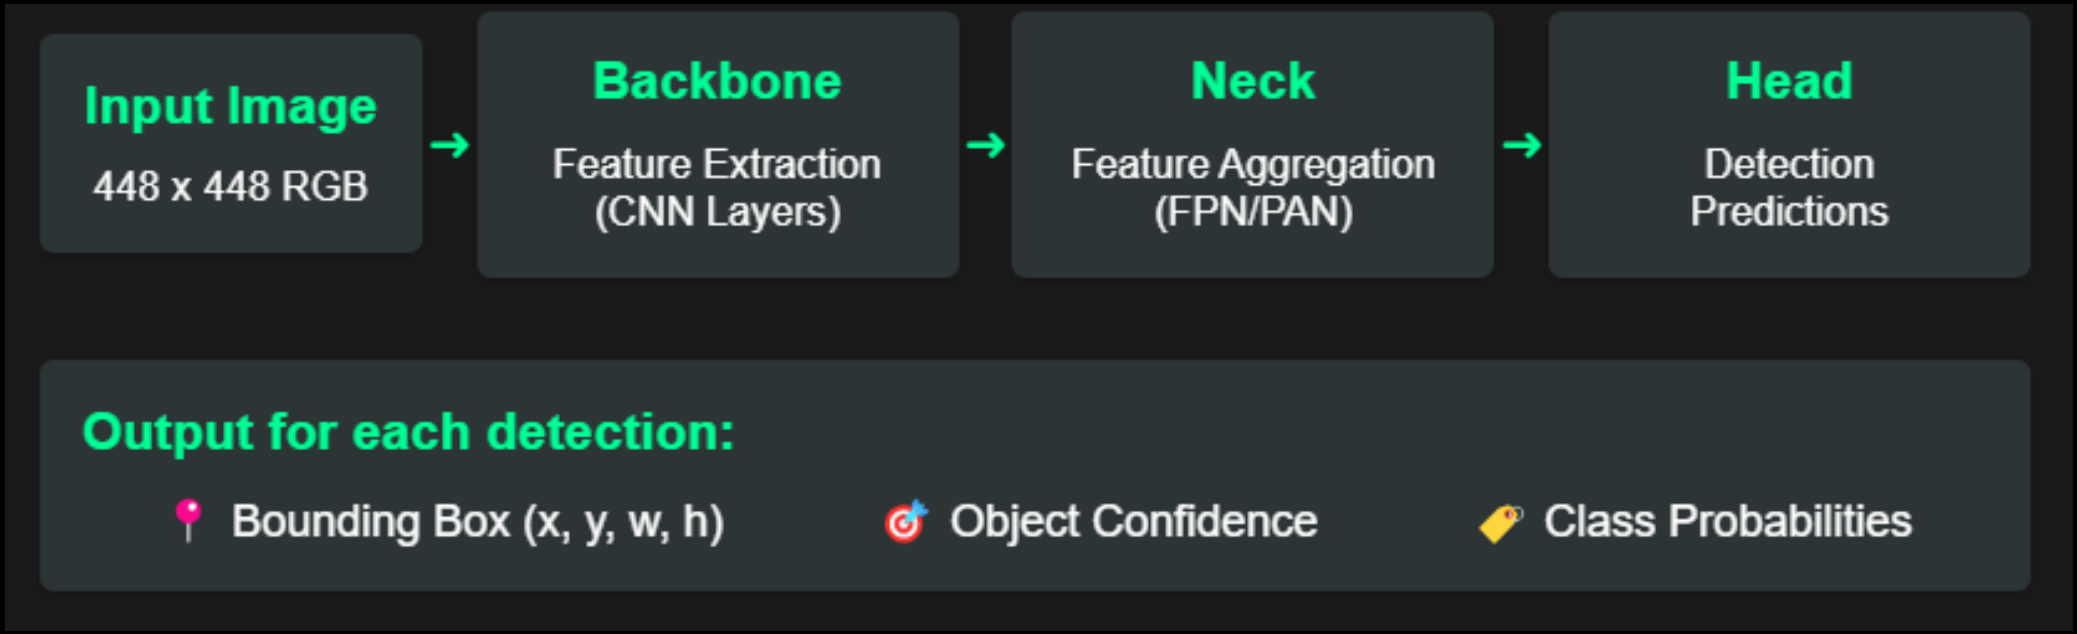

LOSS FUNCTION

Total Loss = $L_1 + L_2 + L_3 + L_4 + L_5$

L_1: Localization Loss

$\lambda_{coord} \Sigma[(x_{pred} - x_{true})^2 + (y_{pred} - y_{true})^2]$

Penalizes incorrect center positions (x,y)

L_2: Size Loss

$\lambda_{coord} \Sigma[(\sqrt{w}_{pred} - \sqrt{w}_{true})^2 + (\sqrt{h}_{pred} - \sqrt{h}_{true})^2]$

Penalizes incorrect box dimensions (w,h)

L_3: Confidence Loss

$\Sigma(C_{pred} - C_{true})^2$

Penalizes incorrect object confidence

L_4: No-object Loss

$\lambda_{noobj} \Sigma(C_{pred} - C_{true})^2$

Penalizes false positive detections

L_5: Classification Loss

$\Sigma \Sigma(p_{pred}(c) - p_{true}(c))^2$

Penalizes incorrect class predictions

Key Points:

- λ_{coord}: Typically 5, gives more weight to spatial predictions
- λ_{noobj}: Typically 0.5, reduces impact of background
- Square root in size loss helps normalize impact of small vs large boxes

NECK

Feature Flow Process

1. Backbone → generates initial features
2. FPN → top-down pathway (large → small)
3. PAN → bottom-up path (small → large)
4. Final features sent to detection heads

Feature Pyramid Network (FPN)

Purpose: Top-down pathway to build feature pyramids

- Large features → detect large objects
- Small features → detect small objects
- Creates multi-scale feature maps (e.g., P3, P4, P5)
- Uses lateral connections to preserve spatial information

Path Aggregation Network (PAN)

Purpose: Bottom-up path augmentation

- Enhances flow of low-level features
- Shortens information path
- Improves feature hierarchy
- Better propagation of low-level patterns

Non-Max Suppression (NMS)

Post-Processing Step (Not part of neck)

Input:

- Multiple bounding boxes
- Confidence scores
- IoU threshold

Process:

1. Sort boxes by confidence
2. Keep highest confidence
3. Remove overlapping boxes
4. Repeat until done

KEY TECHNICAL IMPROVEMENTS ACROSS VERSIONS:

- Early Versions (v1-v3)
  - v1: Single-stage detection
  - v2: Batch norm & anchor boxes
  - v3: Multi-scale predictions
- Middle Era (v4-v7)
  - v4: CSP & Mosaic augmentation
  - v5: Adaptive anchors
  - v7: Dynamic head & label assignment
- Modern Era (v8-v11)
  - V8: Anchor-free detection
  - v9: Programmable gradients
  - v11: Transformer integration

ANCHOR-BASED VS ANCHOR-FREE APPROACHES

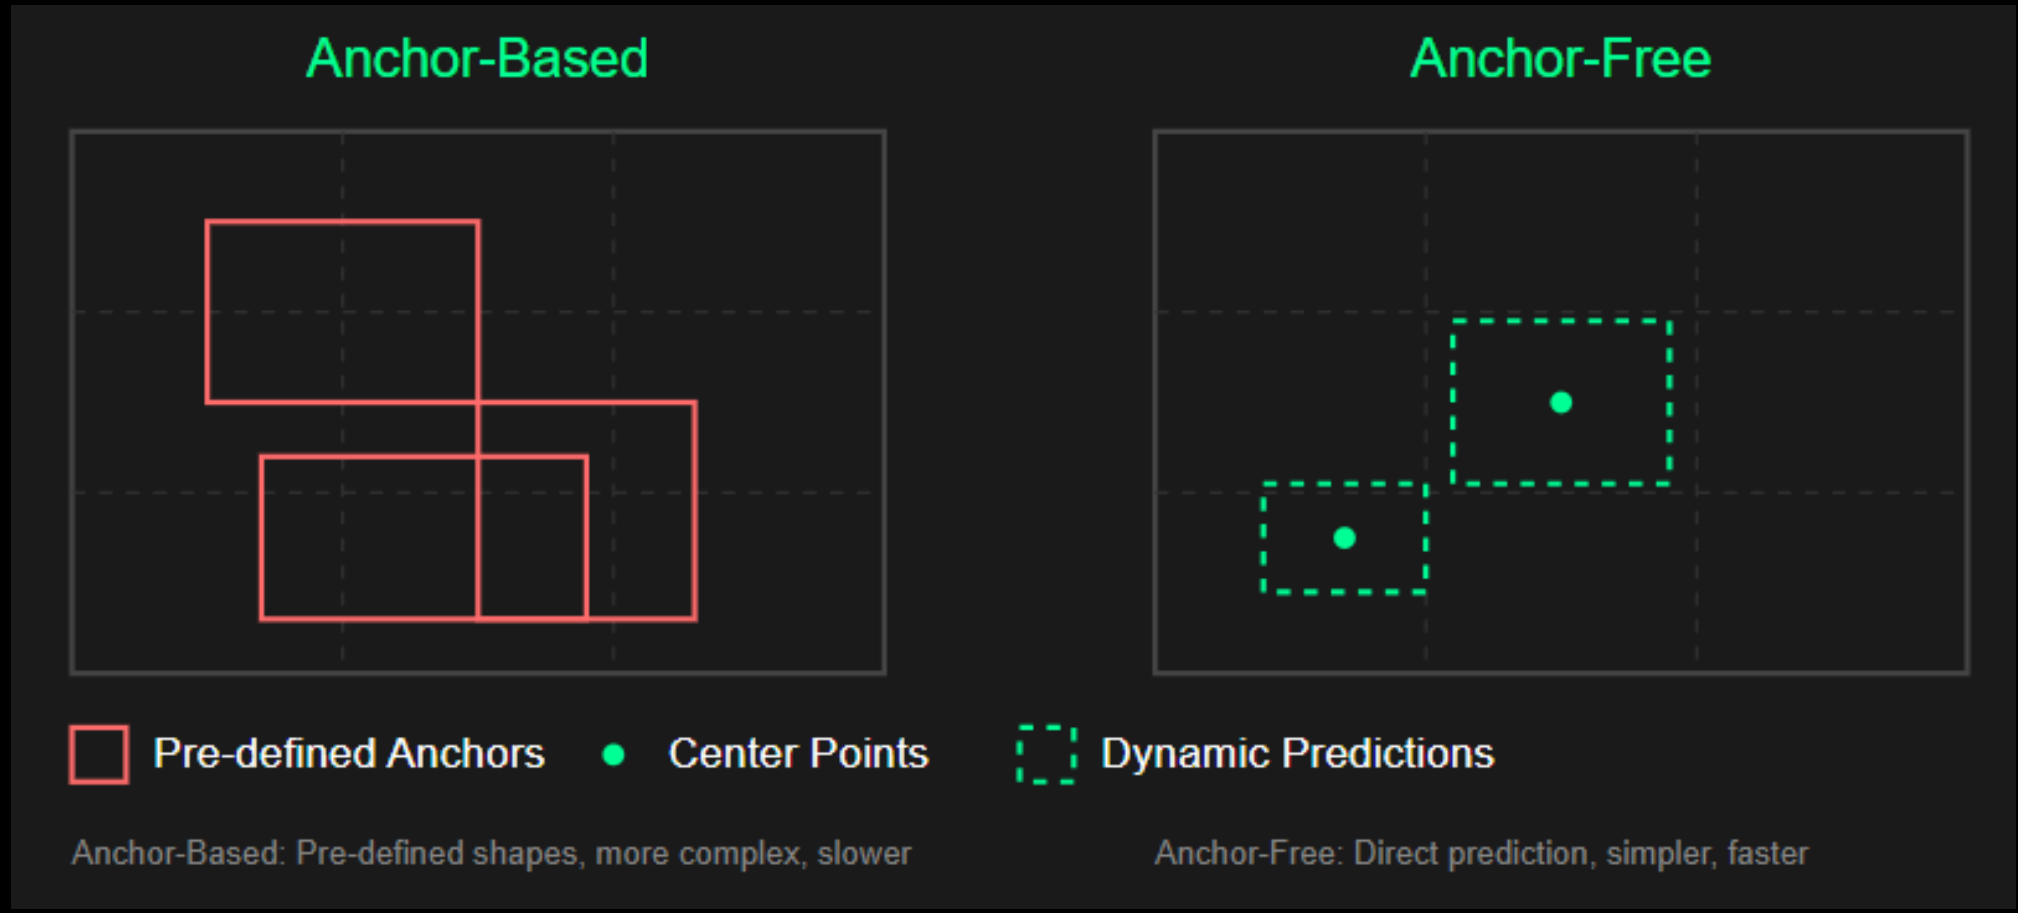

OBJECT TRACKING WITH CNN

- Fully-convolutional network tracker (FCNT)

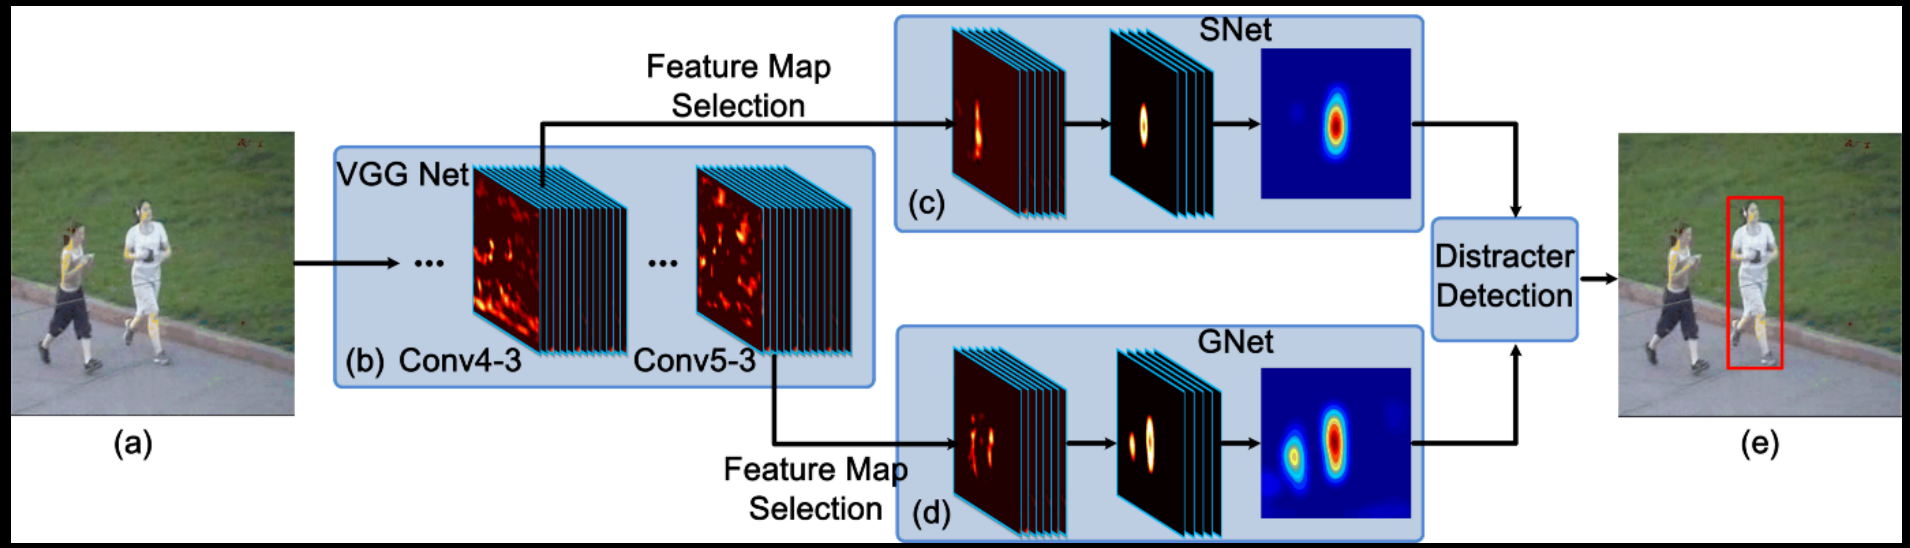

SEMANTIC SEGMENTATION

CLASSIFICATION VS. DETECTION VS. SEMANTIC SEGMENTATION VS INSTANCE SEGMENTATION

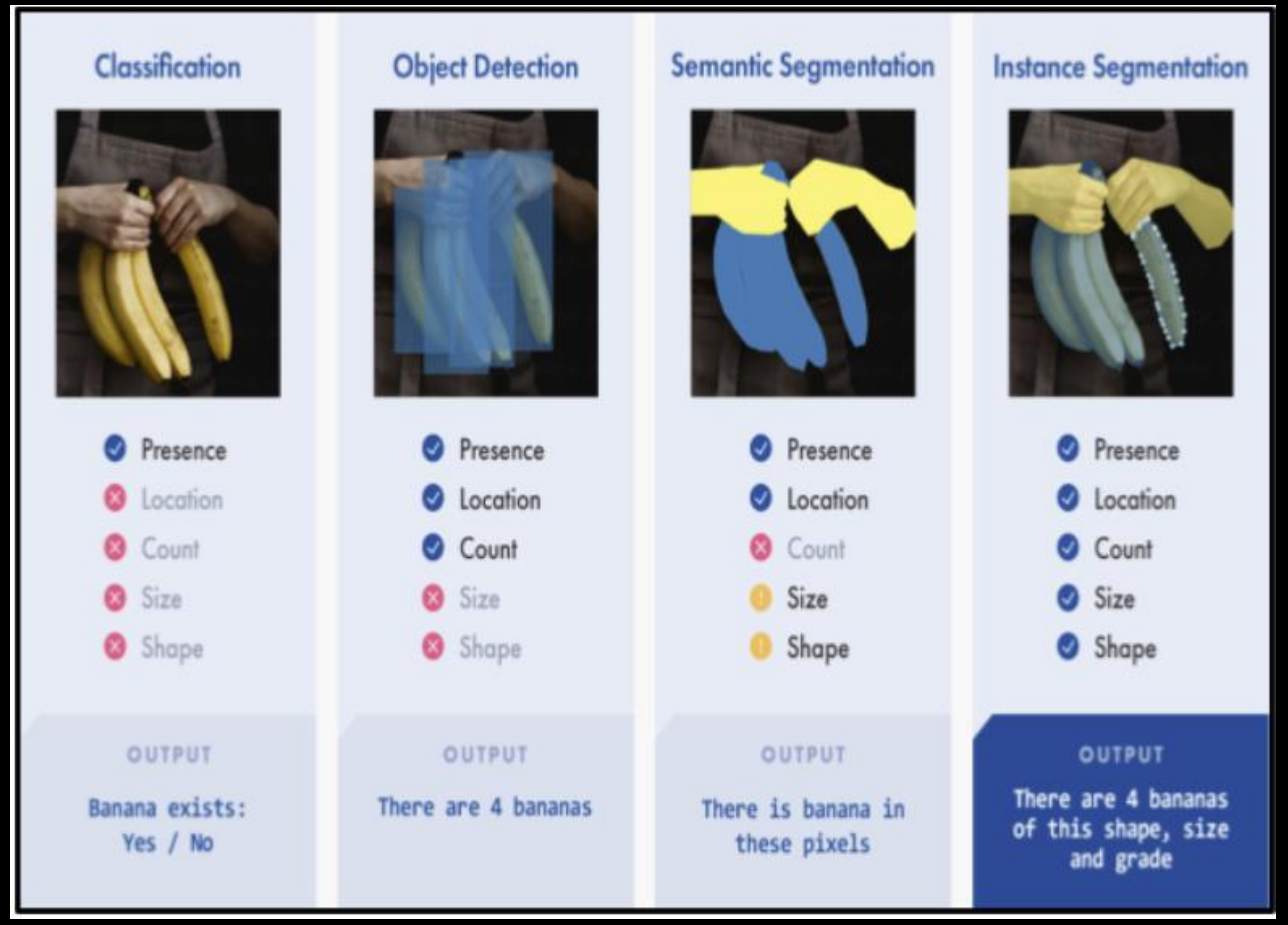

FULLY CONVOLUTIONAL NETWORKS

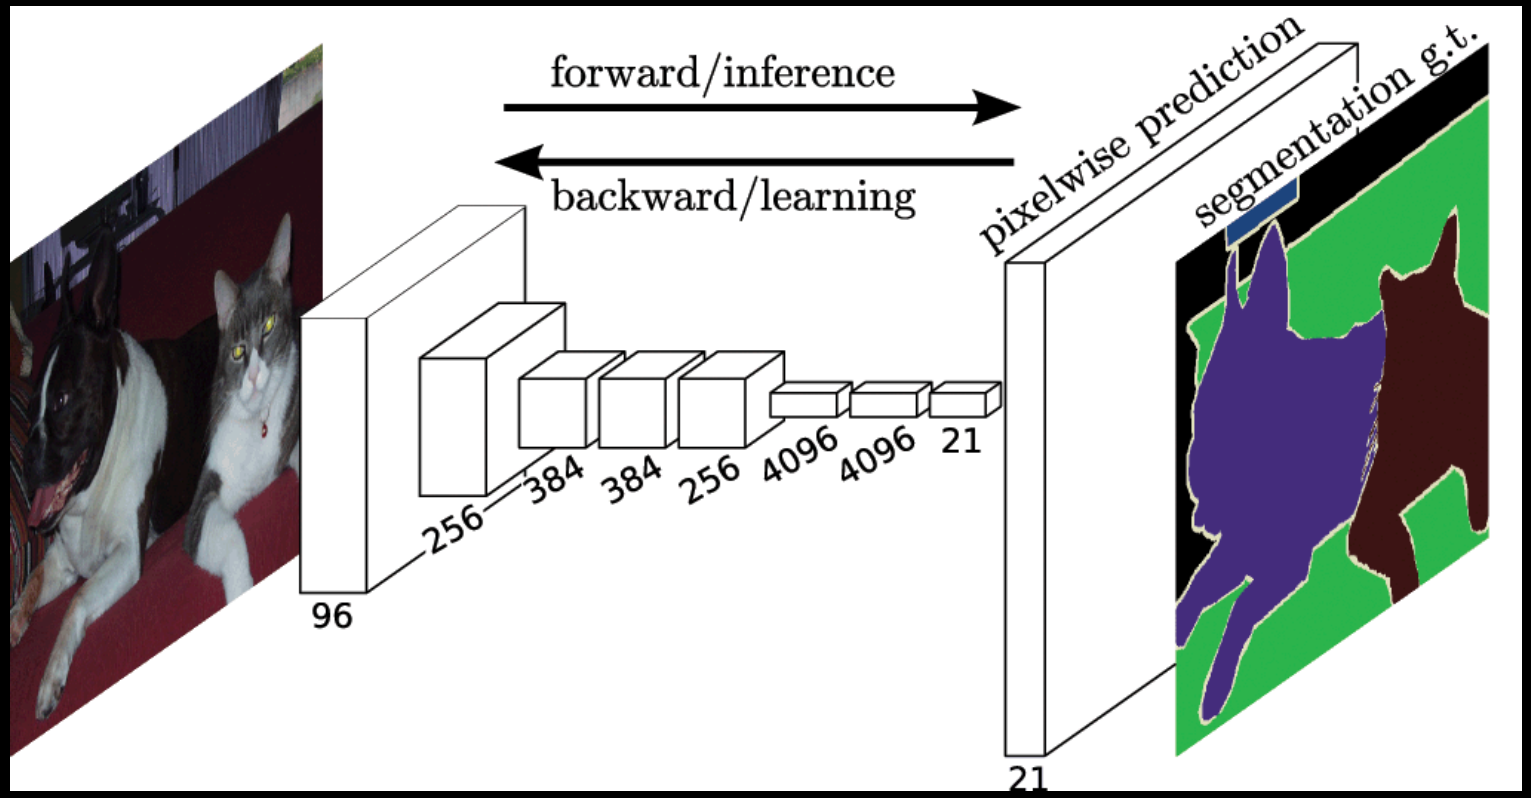

TRANSPOSED CONVOLUTION

- Upsampling layer to recover spatial information

1. stretching
2. filtering

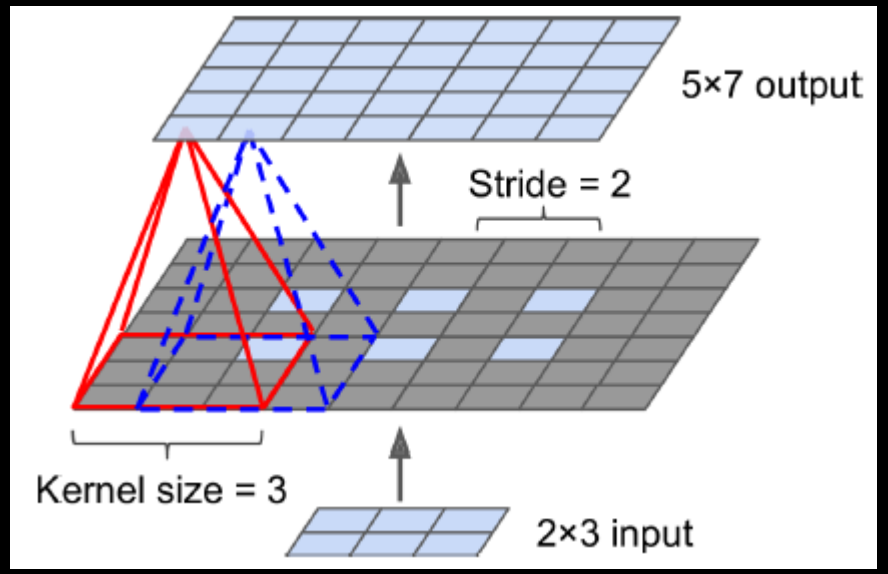

SEGNET

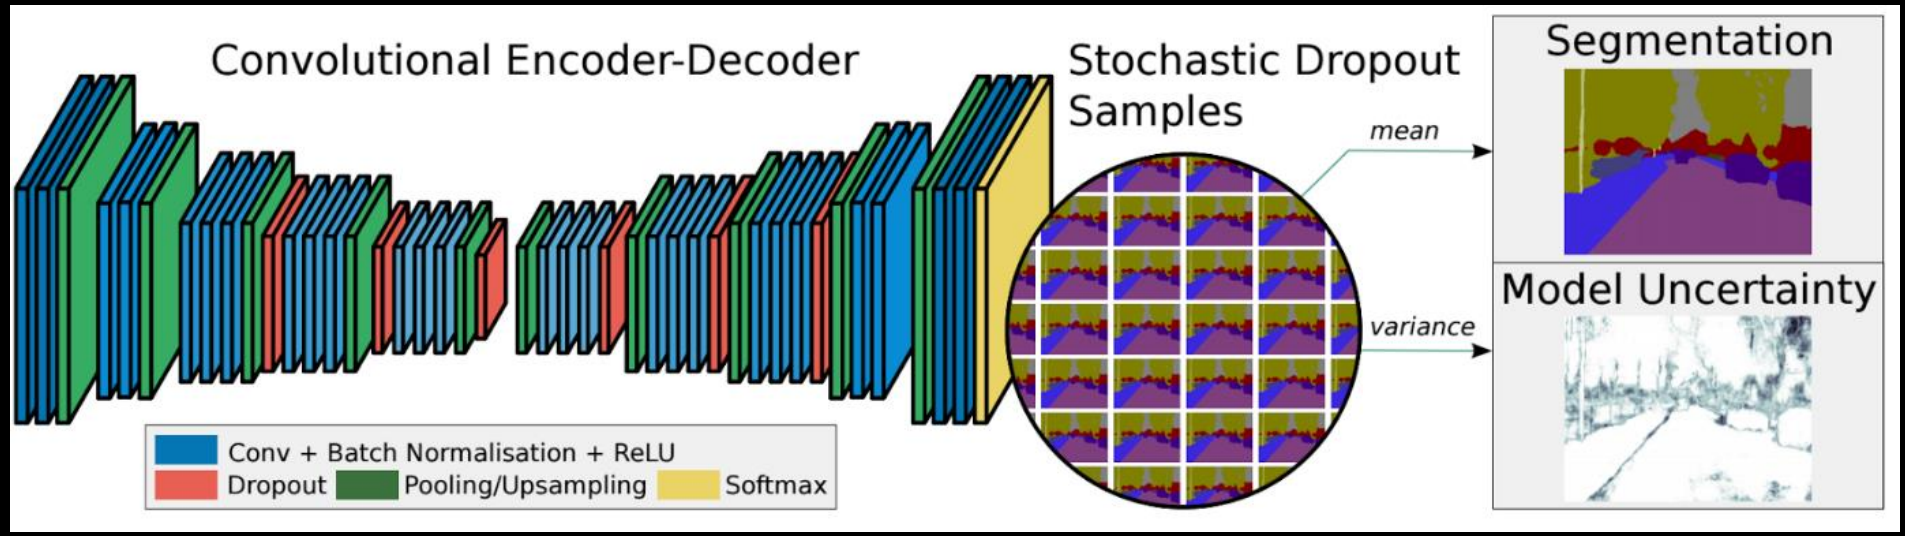
Convolutional Encoder-Decoder
Stochastic Dropout Samples
mean
variance
Segmentation
Model Uncertainty
Conv + Batch Normalisation + ReLU
Dropout
Pooling/Upsampling
Softmax

RECALL DATA AUGMENTATION

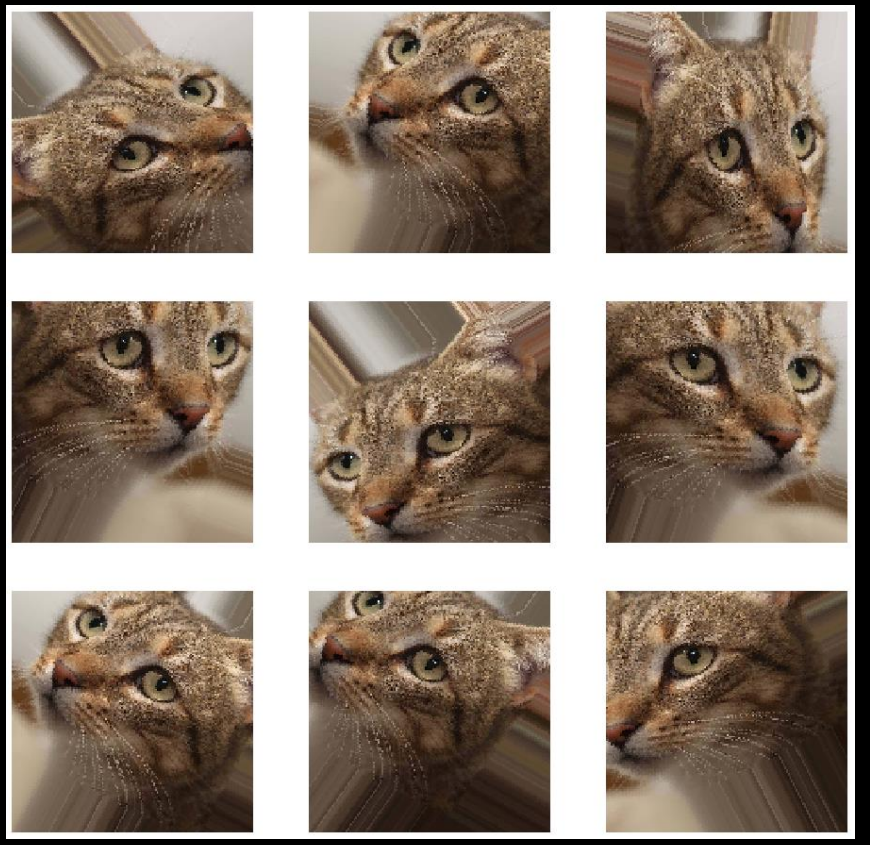

AUTOENCODERS

- Generating augmented data

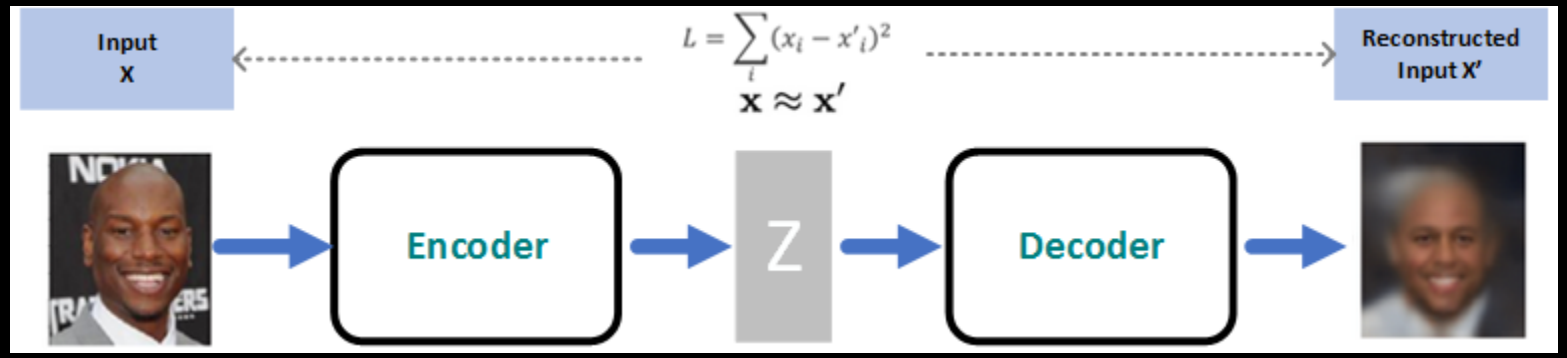

CONDITIONAL VARIATIONAL AUTOENCODER

- Output determined by latent variables, chance and metadata.