

EVML


MORE CNN HANDSON

JEROEN VEEN

AGENDA

- Limitations of neural nets
- Either restart the Keras tuner exercise
- Or build your own CNN
- Or starting with transfer learning


LIMITATIONS OF NEURAL NETS



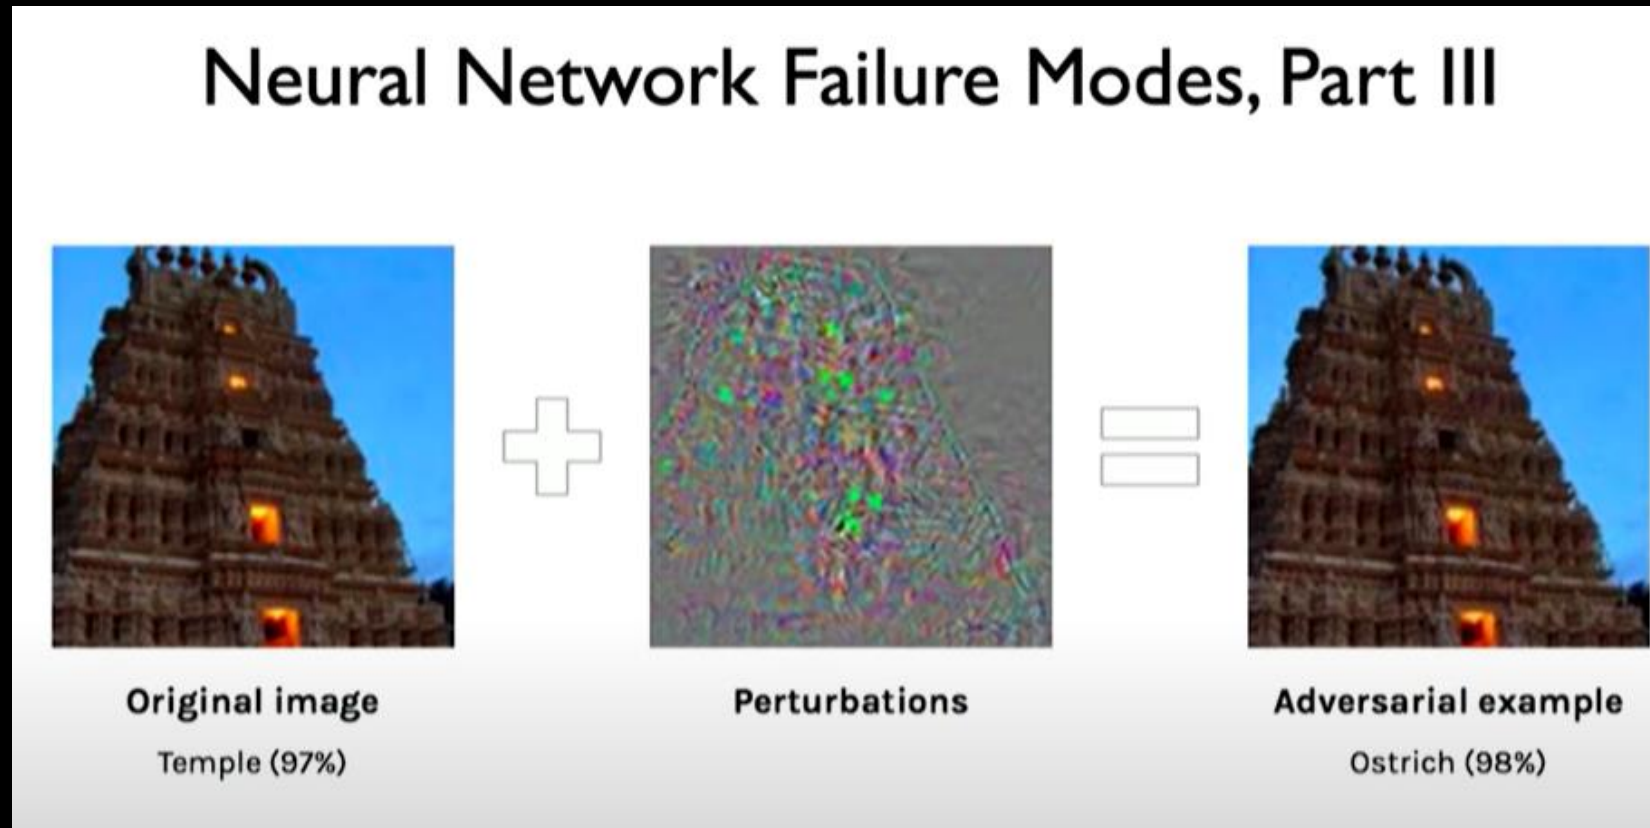
Poor at representing uncertainty, difficult to trust

LIMITATIONS

Neural networks are excellent function approximators
...when they have training data



LIMITATIONS OF NEURAL NETS



Easily fooled by adversarial examples

Source: [MIT Deep Learning 6.S191 \(intro.todeeplearning.com\)](http://intro.todeeplearning.com)

USING THE KERAS TUNER

- Please study:
https://www.tensorflow.org/tutorials/keras/keras_tuner

Use cross-validation to re-evaluate model with the optimal hyperparameters!

- Additional resources:
<https://www.youtube.com/watch?v=O85gh3Ozlul>

BUILD YOUR FIRST CNN

- Train a naive CNN

```
model = keras.Sequential([
    keras.layers.Conv2D(64, 3, activation='relu',
                        input_shape=train_images[0].shape),
    keras.layers.Conv2D(32, 3, activation = 'relu'),
    keras.layers.Flatten(),
    keras.layers.Dense(3, activation='softmax')
])

model.compile(optimizer='adam',
              loss=keras.losses.SparseCategoricalCrossentropy(),
              metrics=['accuracy'])

model.fit(train_images, train_labels, epochs=5, batch_size=32)
```

MORE SOPHISTICATED CNN

```
model = keras.Sequential([
    keras.layers.Conv2D(64,3, activation='relu'),
    keras.layers.MaxPool2D(2,2),
    keras.layers.Dropout(0.5),
    keras.layers.Conv2D(32, 3, activation = 'relu'),
    keras.layers.MaxPool2D(2,2),
    keras.layers.Dropout(0.5),
    keras.layers.Flatten(),
    keras.layers.Dense(64, activation = 'relu'),
    keras.layers.Dense(3, activation='softmax')
])

model.compile(optimizer='adam',
              loss=keras.losses.SparseCategoricalCrossentropy(),
              metrics=['accuracy']
)

model.fit(train_images, train_labels, epochs=5, batch_size=32)
```


RESNET MODEL

- <https://www.kaggle.com/dansbecker/transfer-learning>
- <https://youtu.be/mPFq5KMxKVw>

MOBILENETS

- Efficient Convolutional Neural Networks for Mobile Vision Applications




Figure 3: The Architecture. Our detection network has 24 convolutional layers followed by 2 fully connected layers. Alternating 1×1 convolutional layers reduce the features space from preceding layers. We pretrain the convolutional layers on the ImageNet classification task at half the resolution (224×224 input image) and then double the resolution for detection.

VARIANTS

- V2: M. Sandler et al. 2019, <https://arxiv.org/pdf/1801.04381.pdf>
- MobileNet-Tiny, <https://nitheshsinghsanjay.github.io/>
- Single-Shot Multibox Detector (SSD)

RETRAINING AN IMAGE CLASSIFIER

- https://www.tensorflow.org/hub/tutorials/tf2_image_retraining
- Image size
- Normalization
- Data generator